

# Amazon Sales Analytics – Data Processing & Modeling Synopsis

## 1. Data Exploration & Preparation (Pandas & NumPy)

We use pandas and NumPy to load, inspect, and clean the Amazon sales dataset ([Amazon.csv](#)).

- `pd.read_csv("Amazon.csv")`: Loads the raw sales data into a DataFrame.
- `df.isnull().sum()`: Counts missing values in each column to assess data quality.
- `df.dropna(subset=["OrderDate", "TotalAmount"])`: Removes rows where critical fields (date, target amount) are missing, ensuring reliable analysis.
- `df.nunique(), df["Category"].unique()`: Checks the number of unique values and lists categories, brands, and cities to understand categorical features.
- `df["OrderDate"] = pd.to_datetime(df["OrderDate"])`: Converts the order date from string to proper datetime for time-series analysis.

We also strip column names, remove duplicates, and convert numeric columns (Quantity, UnitPrice, Discount, Tax, ShippingCost, TotalAmount) to numeric types to avoid type errors during modeling.

## 2. Feature Engineering

Several new features are created to make patterns more learnable for models:

- `OrderYear, OrderMonth`: Extracted from `OrderDate` for seasonal and trend analysis.
- `OrderIssue`: Binary target flag indicating whether an order had an issue (Returned / Cancelled = 1, otherwise 0).
- Aggregated customer features (later used for clustering) such as:
  - `NumOrders` (number of orders per customer)
  - `TotalSpend` (sum of `TotalAmount`)
  - `TotalQuantity, AvgDiscount, IssueRate, NumCategories`

These engineered features help the classifier and clustering algorithms capture customer behavior and temporal patterns.

## 3. Feature Transformation & Modeling Setup (Scikit-learn)

We use scikit-learn pipelines to handle preprocessing and modeling in a clean, reproducible way.

### 3.1 Preprocessing

- Numeric features (e.g., Quantity, UnitPrice, Discount, Tax, ShippingCost, TotalAmount, OrderYear, OrderMonth)
  - `SimpleImputer(strategy="mean")`: Fills any remaining numeric missing values.
  - `StandardScaler()`: Normalizes features to zero mean and unit variance so no single variable dominates.
- Categorical features (e.g., Category, Brand, PaymentMethod, City, State, Country)
  - `SimpleImputer(strategy="most_frequent")`: Replaces missing categories with the most common value.
  - `OneHotEncoder(handle_unknown="ignore")`: Converts categories into one-hot encoded vectors suitable for the model.

These transformers are combined in a `ColumnTransformer`, then wrapped in a `Pipeline` with the final estimator.

### 3.2 Core Scikit-learn Flow

- `.fit()`: Learns parameters (means, most frequent categories, encodings) from the training data.
- `.transform()`: Applies these learned parameters to new data.
- `.fit_transform()`: Convenience method used only on training data when fitting and transforming in one step.

## 4. Classification Model – XGBoost

To predict whether an order will have an issue (`OrderIssue`), we train an XGBoost classifier inside the pipeline:

- Model: `XGBClassifier` with tuned hyperparameters (number of estimators, learning rate, max depth, etc.).
- Data split: `train_test_split` with 80% training, 20% test, stratified on `OrderIssue`.
- Metrics: accuracy, precision, recall, F1-score, and full classification report.

This model helps identify risk factors (products, locations, payment types) associated with returns and cancellations.

## 5. Customer Segmentation – KMeans Clustering

On the aggregated customer table:

- Standardize features with `StandardScaler`.
- Try different cluster counts (e.g.,  $k = 2-4$ ) and compute the silhouette score to choose the best  $k$ .
- Assign each customer to a segment (`Cluster`) and summarize profiles:
  - Average spend, number of orders, issue rate, categories purchased.

These clusters reveal groups like “high-value low-issue customers” vs “low-spend high-issue customers,” useful for marketing and support strategies.

## 6. Time-Series Forecasting – ARIMA

Using monthly total sales:

- Build a monthly series: `ts = df.set_index("OrderDate").resample("M") ["TotalAmount"].sum()`.
- Split into train and test (last few months as test).
- Fit an ARIMA model on the training period.
- Forecast over the test horizon and compute RMSE to evaluate accuracy.
- Generate future forecasts (next N months) to estimate upcoming revenue trends.

This shows seasonality and growth patterns in sales.

## 7. NLP-Based Similar Product Recommender

We construct a content-based recommender using product text:

- Combine `ProductName`, `Category`, and `Brand` into a single text field.
- Clean text (lowercasing, removing special characters) for consistency.
- Use `TfidfVectorizer` (uni-grams and bi-grams) to create a TF-IDF matrix.
- Compute cosine similarity between products to find the most similar items.

A helper function `similar_items(product_name, top_n)` returns top-N recommended products based on this similarity, enabling “customers also viewed” style suggestions.

## 8. Saving Artifacts & Deployment Interfaces

To make the project reusable and deployable:

- Save trained models and artifacts with `pickle`:
  - `orderissue_model.pkl` (pipeline with XGBoost)
  - `customer_clusters.pkl` (cluster assignments / centroids)
  - `arima_sales_model.pkl` (time-series model)
  - `recommender.pkl` (products, index mapping, similarity matrix)

## Desktop GUI (Tkinter)

A simple Tkinter-based dashboard provides:

- Buttons to show:
  - EDA charts (payment method distribution, sales by category).
  - Time-series forecast plot (train, test, forecast).
- An input box to type a product name and get similar item recommendations in the UI.

## Flask API

A lightweight Flask API exposes:

- `/predict_order_issue` (POST): JSON input of order features → predicted issue flag + probability.
- `/recommend` (GET/POST): product name → list of similar products.

These interfaces demonstrate how the analytics can be integrated into applications or dashboards.